

databricksPowering the Road_ A comprehensive Analysis on Electric Vehicles (1)

(<https://databricks.com>)

```
#read dataset
df1 = spark.read.format("csv").option("header", "true").load("dbfs:/FileStore/shared_uploads/mgrover3@gmu.edu/Electric Ve
```

```
#display the dataframe
display(df1)
```

Table

	VIN (1-10)	County	City	State	Postal Code	Model Year	Make
1	5UXKTC059G	Yakima	Zillah	WA	98953	2016	BMW
2	5YJ3E1EA2J	Snohomish	Edmonds	WA	98020	2018	TESLA
3	1G1RE6E4XE	Kitsap	Port Orchard	WA	98367	2014	CHEVROLET
4	2C4RC1L76M	Skagit	Bow	WA	98232	2021	CHRYSLER
5	5YJ3E1EA2J	Thurston	Olympia	WA	98513	2018	TESLA
6	WA1E2BFY8N	Snohomish	Snohomish	WA	98296	2022	AUDI

7,404 rows | Truncated data

```
#display summary for the dataframe.
display(df1.summary())
```

Table

	summary ▲	VIN (1-10) ▲	County ▲	City ▲	State ▲	Postal Code ▲	Model Year ▲	Mileage ▲
1	count	173533	173528	173528	173533	173528	173533	173528
2	mean		null	null	null	98174.74609861233	2020.4353523537309	173528
3	stddev	null	null	null	null	2411.1096851358934	2.99444160055937	null
4	min	1C4JJXN60P	Ada	Aberdeen	AE	01545	1997	AL
5	25%	null	null	null	null	98052.0	2018.0	null
6	50%	null	null	null	null	98122.0	2022.0	null

8 rows

```
from pyspark.sql.functions import when, count, col, mean
```

```
#count number of null values in each column of DataFrame
df1.select([count(when(col(c).isNull(), c)).alias(c) for c in df1.columns]).show()
```

VIN (1-10)	County	City	State	Postal Code	Model Year	Make	Model	Electric Vehicle Type	Clean Alternative Fuel Vehicle (CAFV) Eligibility	Electric Range	Base MSRP	Legislative District	DOL Vehicle ID	Vehicle Location	Electric Utility	2020 Census Tract	
0	0	5	5	0	5	0	0	0						10	0	5	5

```
# Calculate mode of the categorical column
mode_value = df1.groupBy("County").count().orderBy(col("count").desc()).select("County").limit(1).collect()[0][0]

# Replace null values with mode
imputed_df = df1.withColumn("County", when(col("County").isNull(), mode_value).otherwise(col("County")))
mode_value = df1.groupBy("City").count().orderBy(col("count").desc()).select("City").limit(1).collect()[0][0]
imputed_df = imputed_df.withColumn("City", when(col("City").isNull(), mode_value).otherwise(col("City")))
mode_value = df1.groupBy("Postal Code").count().orderBy(col("count").desc()).select("Postal Code").limit(1).collect()[0][0]
imputed_df = imputed_df.withColumn("Postal Code", when(col("Postal Code").isNull(), mode_value).otherwise(col("Postal Cod
mode_value = df1.groupBy("Legislative District").count().orderBy(col("count").desc()).select("Legislative District").limi
imputed_df = imputed_df.withColumn("Legislative District", when(col("Legislative District").isNull(), mode_value).otherwi

# Calculate mean of the numeric column
mean_value = df1.agg(mean(col("Electric Range"))).collect()[0][0]

# Replace null values with mean
imputed_df = imputed_df.withColumn("Electric Range", when(col("Electric Range").isNull(), mean_value).otherwise(col("Elec

# Show the DataFrame with imputed values
display(imputed_df)
```

Table

	VIN (1-10)	County	City	State	Postal Code	Model Year	Make
1	5UXKT0C59G	Yakima	Zillah	WA	98953	2016	BMW
2	5YJ3E1EA2J	Snohomish	Edmonds	WA	98020	2018	TESLA
3	1G1RE6E4XE	Kitsap	Port Orchard	WA	98367	2014	CHEVROLET
4	2C4RC1L76M	Skagit	Bow	WA	98232	2021	CHRYSLER
5	5YJ3E1EA2J	Thurston	Olympia	WA	98513	2018	TESLA
6	WA1E2BFY8N	Snohomish	Snohomish	WA	98296	2022	AUDI

7,404 rows | Truncated data

```
#count number of null values in each column of DataFrame
imputed_df.select([count(when(col(c).isNull(), c)).alias(c) for c in imputed_df.columns]).show()
```

-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+ -----+ VIN (1-10) County City State Postal Code Model Year Make Model Electric Vehicle Type Clean Alternative Fuel Vehicle (CAF V) Eligibility Electric Range Base MSRP Legislative District DOL Vehicle ID Vehicle Location Electric Utility 2020 Census Tract -----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+ -----+ 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 10 5 5 -----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+ -----+												
--	--	--	--	--	--	--	--	--	--	--	--	--

```
# Count distinct values in the Vehicle Location column
len(imputed_df.select('Vehicle Location').distinct().collect())
```

Out[31]: 847

```
# Filter rows where 'Vehicle Location' or 'Base MSRP' is null

imputed_df = imputed_df.filter(imputed_df['Vehicle Location'].isNotNull())
imputed_df = imputed_df.filter(imputed_df['Base MSRP'].isNotNull())
```

#check for null values in all columns after filtering


```
imputed_df.select([count(when(col(c).isNull(), c)).alias(c) for c in imputed_df.columns]).show()
```

+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+ +-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+ -----+ VIN (1-10) County City State Postal Code Model Year Make Model Electric Vehicle Type Clean Alternative Fuel Vehicle (CAF V) Eligibility Electric Range Base MSRP Legislative District DOL Vehicle ID Vehicle Location Electric Utility 2020 Census Tract +-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+ -----+ 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 +-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+ -----+													
---	--	--	--	--	--	--	--	--	--	--	--	--	--

```
# imputed_df.(df.State < 40).count()
df_without_WA = imputed_df.filter(imputed_df.State != 'WA')
display(df_without_WA.groupBy('State').count())
```

Table

Visualization 1



42 rows

#filter the dataframe to include records from the WA state.

```
df_with_WA = imputed_df.filter(imputed_df.State == 'WA')
display(df_with_WA)
```

Table

	VIN (1-10)	County	City	State	Postal Code	Model Year	Make
1	5UXKT0C59G	Yakima	Zillah	WA	98953	2016	BMW
2	5YJ3E1EA2J	Snohomish	Edmonds	WA	98020	2018	TESLA
3	1G1RE6E4XE	Kitsap	Port Orchard	WA	98367	2014	CHEVROLET
4	2C4RC1L76M	Skagit	Bow	WA	98232	2021	CHRYSLER
5	5YJ3E1EA2J	Thurston	Olympia	WA	98513	2018	TESLA
6	WA1E2BFY8N	Snohomish	Snohomish	WA	98296	2022	AUDI

7,404 rows | Truncated data

```
#count the number of records for each city.

display(df_with_WA.groupBy('City').count())
```

Table


	City	count
1	Bingen	7
2	Edmonds	1850
3	Bow	94
4	Pasco	537
5	Tumwater	618
6	Auburn	1854

464 rows

```
from pyspark.sql import functions as F

#group by city and count the number of records for each city and filter cities which have mor ethan 900 records.
df_WA_city_count = df_with_WA.groupBy('City').count()
df_WA_city_count_more900 = df_WA_city_count.filter(F.col('count') > 900)
display(df_WA_city_count_more900.orderBy('count'))
```

Table Visualization 1





39 rows

```
brand_w_count = imputed_df.groupBy('Make').count()
brand_with_more_customers = brand_w_count.filter(F.col('count') > 5000)
brand_with_less_customers = brand_w_count.filter((F.col('count') < 5000) & (F.col('count') > 1000))

display(brand_with_more_customers)
```

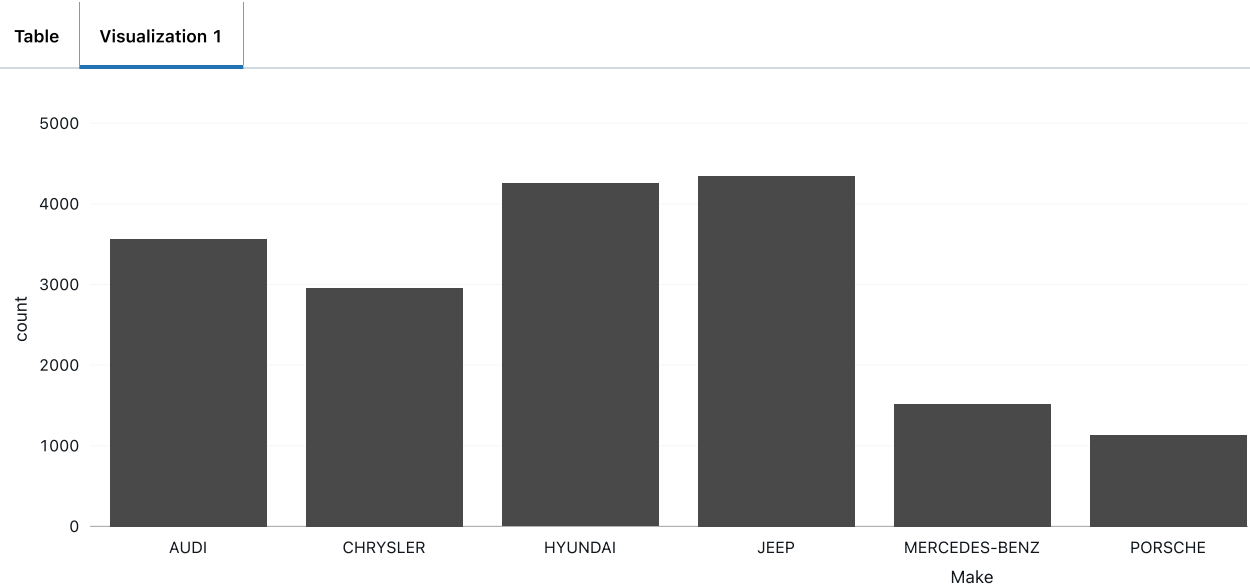
Table Visualization 1






7 rows

```
display(brand_with_less_customers)
```



9 rows

```
df_brand_range = imputed_df.groupBy('Make').agg({'Electric Range': 'mean'})
df_brand_range = df_brand_range.filter(F.col('avg(Electric Range)') > 20)
display(df_brand_range)
```




Make

25 rows

```
year_count = imputed_df.groupBy('Model Year').count()
year_count = year_count.filter(F.col('count')>1000)
display(year_count)
```

Table Visualization 1



13 rows

```
#list numeric columns and select it from dataframe
numeric_column = ['Electric Range','Base MSRP','Legislative District']
numeric_column_df = df_with_WA.select(numeric_column)
display(numeric_column_df)
```

Table

	Electric Range ▲	Base MSRP ▲	Legislative District ▲
1	14	0	15
2	215	0	21
3	38	0	26
4	32	0	40
5	215	0	2
6	23	0	1

10,000 rows | Truncated data

```
#Convert numeric column to int.
for column in numeric_column:
    df_with_WA = df_with_WA.withColumn(column, col(column).cast('int'))

#Split dataframe into test and train data
trainDF, testDF = df_with_WA.randomSplit([0.8, 0.2], seed=42)

print(trainDF.cache().count())
print(testDF.count())
```

138714
34437

```

import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

from pyspark.ml.feature import StringIndexer, OneHotEncoder
from pyspark.ml.classification import RandomForestClassifier
from pyspark.ml.feature import VectorAssembler

# Define categorical column
categoricalCols = [ 'County', 'City', 'State', 'Make', 'Model', 'Electric Vehicle Type', 'Vehicle Location', 'Electric Ut
categoricalCols.remove('State')
# Initialize StringIndexer for categorical features
stringIndexer = StringIndexer(inputCols=categoricalCols, outputCols=[x + "Index" for x in categoricalCols],handleInvalid=

# Initialize OneHotEncoder for the output of StringIndexer
encoder = OneHotEncoder(inputCols=stringIndexer.getOutputCols(), outputCols=[x + "OHE" for x in categoricalCols])

# Initialize StringIndexer for the label/target column
labelToIndex = StringIndexer(inputCol="Clean Alternative Fuel Vehicle (CAFV) Eligibility", outputCol="label")

# Combine all feature columns into a single feature vector
assemblerInputs = [c + "OHE" for c in categoricalCols] + numeric_column # Note: using OHE columns
vecAssembler = VectorAssembler(inputCols=assemblerInputs, outputCol="features")

from pyspark.ml import Pipeline
# Define the pipeline based on the stages created in previous steps.
pipeline = Pipeline(stages=[stringIndexer,encoder,vecAssembler,labelToIndex])

# Define the pipeline model.
pipelineModel = pipeline.fit(trainDF)
predDF = pipelineModel.transform(trainDF)
predDF1 = pipelineModel.transform(testDF)
# Define RandomForestClassifier
from pyspark.ml.classification import RandomForestClassifier
# Initialize the RandomForestClassifier
rfClassifier = RandomForestClassifier(featuresCol="features", labelCol="label", numTrees=10)

# Train the model on the transformed DataFrame
rfModel = rfClassifier.fit(predDF)

# Make predictions on the training data (you can also use a separate test set)
predictions = rfModel.transform(predDF1)

```

```
predictions.select("features", "label", "prediction", "probability").show()
```

features	label	prediction	probability
(1290,[0,50,506,5...]	2.0	2.0	[0.20689909026626...
(1290,[0,44,506,5...]	2.0	2.0	[0.20689909026626...
(1290,[0,39,506,5...]	2.0	2.0	[0.20689909026626...
(1290,[0,50,506,5...]	2.0	2.0	[0.20689909026626...
(1290,[0,65,506,5...]	2.0	2.0	[0.20689909026626...
(1290,[0,40,506,5...]	2.0	2.0	[0.20689909026626...
(1290,[0,50,506,5...]	2.0	2.0	[0.20689909026626...
(1290,[0,39,506,5...]	2.0	2.0	[0.20689909026626...
(1290,[0,50,506,5...]	2.0	2.0	[0.20689909026626...
(1290,[0,50,506,5...]	2.0	2.0	[0.20689909026626...
(1290,[0,50,506,5...]	2.0	2.0	[0.20689909026626...
(1290,[1,43,506,5...]	2.0	2.0	[0.21229847726617...
(1290,[3,42,506,5...]	2.0	2.0	[0.20689909026626...
(1290,[3,101,506,...]	2.0	2.0	[0.20689909026626...
(1290,[16,159,506...]	2.0	2.0	[0.20689909026626...

```
| (1290, [0,40,506,5...| 2.0| 2.0| [0.20689909026626...|
| (1290, [0,50,506,5...| 2.0| 2.0| [0.20689909026626...|
```

```
from pyspark.ml.evaluation import MulticlassClassificationEvaluator

# Calculate the accuracy of the predictions

evaluator = MulticlassClassificationEvaluator(labelCol="label", predictionCol="prediction", metricName="accuracy")
accuracy = evaluator.evaluate(predictions)

print("Accuracy:", accuracy)
Accuracy: 0.9006010976565902
```

```
#Extract the feature columns from the DataFrame and retrieve feature importances.

feature_columns = predDF.columns[:-1]
feature_importances = rfModel.featureImportances

# Create a dictionary of feature importances
importances_dict = {feature_columns[i]: feature_importances[i] for i in range(len(feature_columns))}

# Print feature importances
for feature, importance in importances_dict.items():
    if importance >0:
        print(f"Feature: {feature}, Importance: {importance}")
Feature: Vehicle LocationIndex, Importance: 4.08680900413647e-05
Feature: Electric UtilityIndex, Importance: 3.5202241936741704e-05
Feature: ModelOHE, Importance: 3.818257512498391e-05
```

```
#reaserach question 2(Which city consistently has electric cars that can drive the farthest on a single charge, considerin
from pyspark.sql.functions import col, min,max ,avg
from pyspark.sql.functions import desc, asc

# Average of electric range of each car model in each city.
city_avg_range = df_with_WA.groupBy("City", "Model Year", "Make", "Model","State").agg(avg("Electric Range").alias("Avg C
Electric Range In City"))

#Average of each car model.
car_avg_range = df_with_WA.groupBy("Model Year", "Make", "Model").agg(avg("Electric Range").alias("Avg Car Electric Range

# Perform an inner join between the DataFrames city_avg_range and car_avg_range.

inner_join_df = city_avg_range.join(car_avg_range, ["Model Year", "Make", "Model"], "inner")
display(inner_join_df.orderBy(desc("Avg Car Electric Range In City")))
```

Table						
	Model Year ▲	Make ▲	Model ▲	City ▲	State ▲	Avg Car Ele
1	2020	TESLA	MODEL S	Newcastle	WA	337
2	2020	TESLA	MODEL S	Tukwila	WA	337
3	2020	TESLA	MODEL S	Battle Ground	WA	337
4	2020	TESLA	MODEL S	Black Diamond	WA	337
5	2020	TESLA	MODEL S	Chattaroy	WA	337
6	2020	TESLA	MODEL S	Woodland	WA	337

10,000 rows | Truncated data


```
# Filter DataFrame to include only records where the average electric range in the city is greater than the overall avera

filter_range = inner_join_df.filter(inner_join_df["Avg Car Electric Range In City"] > inner_join_df["Avg Car Electric Ran
display(filter_range.orderBy(desc("Avg Car Electric Range In City")))
```


Table

	Model Year	Make	Model	City	State	Avg Car Electric Range In City
1	2020	TESLA	MODEL S	Battle Ground	WA	337
2	2020	TESLA	MODEL S	Newcastle	WA	337
3	2020	TESLA	MODEL S	Tukwila	WA	337
4	2020	TESLA	MODEL S	Chattaroy	WA	337
5	2020	TESLA	MODEL S	Woodland	WA	337
6	2020	TESLA	MODEL S	Carnation	WA	337

1,427 rows



18 rows